

CareerArth AI — Executive Summary

Aryan Bhandari · Entrepreneurship Capstone · Submitted to Professor Tushar Jaruhar

Opportunity. Mid-career professionals (5–10 years in) who have quietly stalled face a gap between vague worry and expensive commitment: ChatGPT gives unstructured advice, a coach demands a big engagement before the problem is even named. CareerArth — an existing career-advisory brand whose site already runs a free-assessment → paid-consultation funnel — can own the structured first step in between. (*ICP and opportunity: reasoned hypotheses, not yet customer-validated.*)

Solution. CareerArth AI turns a 5–7 minute structured intake into a deterministic **ARTH scorecard** — Alignment, Risk Exposure, Trajectory, Human Capital (a prototype diagnostic framework, not a validated psychometric instrument) — plus an AI-drafted, evidence-cited analysis grounded in a curated 18-role career knowledge base: matched career paths, skill gaps, a 90-day plan, and a consultation call-to-action. Three design rules protect trust: the AI never computes scores, every claim must cite its evidence (enforced by an automated lint), and a human consultant reviews before anything is final.

MVP — demonstrated. A working branded web product: intake wizard → scorecard → report → consultant review screen → funnel instrumentation. **42/42 automated tests pass.** Six synthetic personas produce genuinely differentiated results (overall scores 42–88, four interpretation bands, distinct career-path sets, clean handling of a career-break profile). *Boundary: these personas are QA fixtures, not customer evidence — and every report to date uses a deterministic mock generator; the Claude integration is fully implemented but has never run live (no API key), so real AI report quality is unobserved.*

Business model — hypothesis. Free diagnostic as lead generation; the paid human consultation (hypothesised ₹1,999–₹4,999) is the revenue event — formalising the motion the production site already runs. Modelling shows estimated AI cost (\$0.05–\$0.15/report, estimated, never measured) is a rounding error; **consultant time and free→paid conversion dominate the economics** (base-case contribution ≈ ₹2,089/customer at an assumed 10% conversion — every input an explicit assumption).

Evidence line. Demonstrated: technical feasibility of the full workflow. Implemented-but-unvalidated: live Claude generation and its cost. Untested: demand, conversion, willingness to pay, pricing, product-market fit. No real customer has used the product; nothing here claims otherwise.

Next step. A ready-to-run 10–20-person pilot — recruitment message, consent, surveys, interview guide, WTP questions, and pre-committed success/failure thresholds (≥80% intake completion; trust ≥3.5/5; ≥30% consultation interest; median WTP ≥₹1,500) — requiring **zero new engineering**, because the MVP already instruments every funnel step. Then a 30/60/90-day roadmap ending in an evidence-based **continue / revise / pivot / stop** decision.