

ENTREPRENEURSHIP CAPSTONE

CareerArth AI

A structured career diagnostic that turns quiet stagnation into a named problem — and a booked consultation.

Submitted by Aryan Bhandari

Submitted to Professor Tushar Jaruhar · August 2026

PROBLEM

Careers don't collapse. They erode.

- Mid-career professionals stall quietly — solid reviews, no movement, narrowing options
- They feel it, but can't decompose it: **skills? industry? network? ambition?**
- Their choices are extreme (quit / do nothing) or expensive (hire a coach, unbriefed)

The pain is **epistemic** — they cannot name the problem. Naming it is a product.

Beachhead: the stalled mid-career professional

- **5–10 years in** · no promotion or scope change in 2+ years · currently employed
- Tech, finance, marketing/consulting — the footprint our knowledge base covers
- Triggers: denied promotion · a peer's title change · a tenure milestone

Hypothesis: they won't pay for a quiz — they'll pay for a consultation once the diagnostic proves it understands them. This is the venture's #1 untested assumption.

Everything they try today fails a different way

Option	Failure mode
ChatGPT / generic AI	Fluent, unstructured, ungrounded — quality depends on their prompting
Career coach	Big commitment demanded before the problem is even named
LinkedIn	A feed, not a decision
Personality tests	Describe traits, not this situation
Free content	The user does all the synthesis

The gap isn't "AI career advice" — that's free. It's a **cheap, structured, trustworthy first step**, owned by the brand that sells the next step.

SOLUTION

7 minutes → a scored, evidence-cited diagnostic → a human

- Structured intake — profile + 20-item ARTH questionnaire
- Deterministic ARTH scorecard — **no AI in the numbers**
- AI analysis grounded in a curated career knowledge base — **every claim cites its evidence**
- Career paths with trade-offs · skill gaps · 90-day plan
- **A human consultant reviews before anything is final**

The three gold rules are the trust architecture — and the answer to "why not just ChatGPT?"

ARTH — four questions, one scorecard

Dimension	Question	Weight
Alignment	Does today's work build toward the ambition?	0.30
Risk Exposure	How insulated from industry & automation shifts?	0.25
Trajectory	Are options expanding or narrowing?	0.20
Human Capital	Network, brand, transferability?	0.25

Five bands: Strong 85–100 · Stable 70–84 · Watch 55–69 · Vulnerable 40–54 · Critical 0–39

A prototype structured framework — not a validated psychometric instrument. Disclosed on the intake page, in every score, and on every report.

WORKING MVP — DEMONSTRATED

Built, tested, instrumented


42/42
automated tests pass

42→88
persona score spread

18
curated roles (O*NET)

- Branded flow: intake → scorecard → report → consultant review
- Full funnel instrumentation — pilot-ready

Boundary: all reports to date are mock-generated (Claude integration implemented, never run live) · personas are QA, not customers



Free diagnostic → paid consultation

- The diagnostic is **lead generation**, not the product
- Revenue event: the human consultation — hypothesised **₹1,999–₹4,999**
- Formalises the funnel the CareerArth site **already runs** (free assessment → "Talk to an Expert")
- Phase 2 option: premium AI + human-reviewed strategy bundle — *after* conversion evidence

Chosen from three models: it fits the MVP's real delivery order and the trust posture — no payment before value is shown.

The AI is a rounding error. The human is the cost.

Scenario	Price	Assumed conversion	Consultant session	Contribution / customer
Conservative	₹1,999	5%	60 min @ ₹1,500/hr	≈ ₹194
Base	₹2,999	10%	45 min @ ₹1,000/hr	≈ ₹2,089
Upside	₹4,999	20%	30 min @ ₹800/hr	≈ ₹4,453

Estimated AI cost \$0.05–0.15/report — estimated from list prices, never measured live. Every input above is a labelled assumption.

Margin hinges on conversion rate + consultant minutes — both testable in a 10–20-person pilot. Optimise those, not tokens.

No tech moat. A trust-and-distribution bet.

Copyable in weeks

- The retrieval pipeline
- The scoring rubric (heuristic, disclosed)
- The prompt design

Hard to copy

- Existing advisory brand + live funnel
- AI paired with **real** human review
- Evidence-cited discipline as posture

The durable asset must be **built**: funnel data, review learnings, outcome tracking — none exist today, and we say so.

What's proven vs. what's not

Demonstrated	Implemented, unvalidated	Hypothesis
Working MVP · 42/42 tests	Claude integration (never run live)	ICP & demand
Deterministic ARTH scoring	AI cost estimate (\$0.05–0.15)	Conversion & willingness to pay
6 differentiated synthetic personas	Real report quality	Pricing · consultant economics
Funnel instrumentation		Product-market fit

No real customer has used this product. Synthetic personas are QA fixtures — never counted as customer evidence anywhere in this capstone.

The pilot is ready to run — zero new engineering

- **10–20 real users:** consent → diagnostic → report → surveys → WTP → interview
- **Pre-committed thresholds:** ≥80% completion · trust ≥3.5/5 · ≥30% consultation interest · median WTP ≥₹1,500
- **30/60/90 days:** first real Claude reports → pricing anchors + consultant timing → **continue / revise / pivot / stop** memo

Ask: treat this as a live venture experiment with the evidence line drawn honestly — not as a finished claim.