

Appendix A

Business & Validation — ICP, business model, unit economics, funnel, competition, GTM, assumption register, and the ready-to-run pilot kit.

Submitted by Aryan Bhandari

Submitted to Professor Tushar Jaruhar

Course Entrepreneurship Capstone (ENT-CP-4997-1)

Date August 2026

Appendix A — Business & Validation

CareerArth AI · Aryan Bhandari · for Professor Tushar Jaruhar

Consolidates the Stage 3 business analysis and Stage 4 validation kit. Everything in this appendix is a **hypothesis or a designed-but-unexecuted instrument** unless marked [DEMONSTRATED]. Source: docs/stage3/ , docs/stage4/ .

A1. ICP and Job-To-Be-Done

Beachhead: mid-career professionals, 5–10 years of experience, no promotion or material scope change in 2+ years, employed, within the tech/finance/marketing footprint of the 18-role knowledge base.

Why this segment (and not the others considered):

Candidate	Verdict
Early-career (1–3 yrs)	Real pain, diffuse urgency, weakest WTP case — secondary audience
Role/industry switchers	High intent, but their core need (breadth of adjacent-role knowledge) stresses the product's thinnest asset
Stagnating mid-career	Matches the rubric's strongest built-in signal (deterministic tenure penalty on Trajectory) and the production brand's existing "gilded stagnation" narrative

JTBD: "When I've stopped growing but can't tell if the problem is me, my role, or my industry — give me an honest, structured read and a credible next move, without a big commitment first."

Triggers: denied/bypassed promotion · second consecutive "solid" review · peer's title change · tenure milestone.

WTP hypothesis: low for the diagnostic itself; real for the follow-up consultation *after* the diagnostic demonstrates specificity. The single most important untested assumption.

A2. Business model comparison

Criterion	A. Paid AI report		C. Premium bundle
Trust at payment	Low — pay before evidence	High — value first	Medium-high
Scalability	High	Diagnostic scales; revenue capped by consultant hours	Capped
Consultant dependency	Low	High, but only on high-intent clients	Highest
Margins	High if unbundled	Advisory margins; diagnostic = CAC	Highest/unit, services-shaped
Fit with MVP as built	Weak (report ships as draft pre-review)	Strong — mirrors built flow + live site funnel	Needs re-sequencing

Recommendation: B. No product change required; matches the real delivery order; avoids charging for an explicitly-disclaimed heuristic score; C remains the phase-2 evolution once conversion evidence exists.

A3. Pricing & unit economics (all inputs assumptions)

Estimated AI cost/report: **\$0.05–\$0.15 (~₹4–₹13 at an assumed ₹85/\$)** — derived from published list prices (input \$5/MTok, output \$25/MTok) × estimated token volume (~1,500–2,500 in / ~1,500–4,000 out). **Never measured live.**

Contribution = price – (AI cost ÷ conversion) – consultant session – ~2.5% gateway:

Scenario	Price	Conversion	Session	AI amortised	
Conservative	₹1,999	5%	60 min @ ₹1,500/hr	≈ ₹255	≈ ₹194
Base	₹2,999	10%	45 min @ ₹1,000/hr	≈ ₹85	≈ ₹2,089
Upside	₹4,999	20%	30 min @ ₹800/hr	≈ ₹21	≈ ₹4,453

Structural finding: conversion rate and consultant minutes dominate; AI cost is 1–13% of price. Conservative-case margin is thin enough that conversion below ~5% or sessions beyond ~60 min break the model — the pilot's most important numbers.

A4. Funnel map vs. existing instrumentation [instrumentation DEMONSTRATED]

Stage	Desired action	Main friction	Metric	Already measured?
Discovery / landing	Click "Free Diagnostic"	Skepticism, unclear differentiation	landing → intake start	No (production site, out of scope)
Intake	Complete 8-screen wizard	20-item fatigue	intake_completion_rate + per-step drop-off	Yes
Report	Read, engage	Generic feel; trust gap on unvalidated score	report_viewed, usefulness_rating	Yes
Consultation interest	Click CTA	Free-AI → paid-human leap	consultation_interest rate	Yes
Paid service	Book + pay	No booking/payment exists	would need new events	No — known gap
Follow-up/referral	Return, refer	No mechanism	—	No — known gap

A5. Competition (summary)

	Personalisation	Structure	Actionability	Human	Trust	Speed	Price	Career grounding
LLM chat	Med	None	Med	None	Low-Med	Instant	Free	None inherent
Coach	High	Varies	High	Full	High	Slow	High	High
LinkedIn	Low	Low	Low	None	Med	Instant	Free-low	Med
Assessments	Low-Med	High	Low	None	Med	Fast	Free-low	Low
Free content	None	None	Low-Med	None	Low	Instant	Free	Varies
CareerArth	Med-High	High	Med-High	Partial (real review)	Med (earned, not claimed)	Fast	Free → paid consult	Med (18-role curated KB)

Honest moat statement: technology copyable in weeks; the bet is brand + funnel + human-review trust posture; durable assets (funnel data, outcome tracking) must be built.

A6. Lean GTM (first 10–50 users)

1. **Direct outreach** to known ICP-matching contacts (~10) — doubles as pilot recruitment.
2. **Referral ask** inside every pilot conversation (~20).
3. **Honest prototype-tester posts** in career communities (~20). Deferred: alumni networks (wrong ICP fit), partnerships (premature), content (too slow for n=50), paid (no).

A7. Assumption register (top 10, ranked)

#	Assumption	Current evidence	Cheapest test	Threshold
1	Real Claude report >> mock template	None — zero real generations	Generate 3–5, read them	Team judges showable
2	ICP finds report credible/accurate	None	Pilot survey Q1–3	≥3.5/5
3	Free→paid conversion sustains Model B	None (live funnel shape only)	CTA rate (instrumented)	≥30% interest
4	Users complete the 20-item intake	Zero completions ever logged	Completion rate (instrumented)	≥80%, 5–8 min
5	WTP clears viable floor	None	Unanchored WTP question	Median ≥₹1,500
6	Honesty disclaimer doesn't kill trust	None	Survey trust Q + interview	≥3.5/5 despite disclosure
7	Consultant review ≤ ~20 min	One QA point (6.5 min, developer, synthetic)	Time a real review of a real report	≤20 min avg
8	KB covers pilot cohort	Personas were written to fit	<code>currentRoleMatched</code> rate	≥80% matched
9	Production funnel converts non-trivially today	Page copy only	Pull site analytics	Baseline established
10	Users perceive difference vs ChatGPT	None	Interview Q3	Majority name a specific reason

A8. Pilot protocol (designed, not executed)

Prerequisite: generate and read 3–5 real Claude reports before recruiting anyone.

Recruit (15–25 invited → 10–20 complete) → consent → pre-survey (4 Q) → real intake → real report → post-survey (8 Q) → WTP (4 Q) → optional 10-min interview (5–8 people) → CTA-click measurement. Every quantitative step uses existing instrumentation; a one-tab spreadsheet (schema drafted in Stage 4) holds the rest. Real participants: `isSynthetic: false`, recorded consent, data kept out of the repository.

A9. Research instruments (ready to send — full text in `docs/stage4/05`)

- **Recruitment message** (short, honest, "prototype — no sales pitch")
- **Consent text** (AI-generated + human-reviewed, not licensed counselling; local storage; deletion on request; 18+; voluntary)
- **Pre-survey:** prompt, stuck-ness 1–5, prior tools, definition of "useful"
- **Post-survey:** accuracy · specificity · trust (1–5); surprise; decision impact (clearly/somewhat/no); consultation likelihood 1–5; perceived vs actual time
- **WTP block:** unanchored consultation price → diagnostic anchors ₹0/₹299/₹999 → model preference A/B/C → repeat-use intent

- **Interview guide:** walk-through reaction · disagreement moment · tell-a-friend · what would make you book/pay · what we got wrong

A10. Success/failure thresholds (pre-committed)

Metric	Success	Failure
Intake completion	≥ 80%	< 50%
Median intake time	5–8 min	> 12 min
Accuracy/trust (avg)	≥ 3.5/5	< 3.0/5
Decision-clarifying effect	≥ 50% clearly/somewhat	< 25%
Consultation interest	≥ 30%	< 10%
WTP (median, unanchored)	≥ ₹1,500	< ₹500 / majority "wouldn't pay"
Preferred model	≥ 50% for one option	3-way split
Consultant review	≤ 20 min avg	> 40 min / heavy rewrites

At n=10–20, any metric within ~10% of threshold is inconclusive; interview evidence governs. Results feed the day-90 continue/revise/pivot/stop memo.